

# Learning from Data – Assignment 2: Neural Networks

## General remarks

This assignment has to be completed **individually**. During the lab, we will have a **small live competition!** See below for details. The best performing students get a bonus on their grade.

This assignment is meant to get you acquainted with the basics of neural networks: by implementing an MLP and experimenting with word embeddings. First, you are provided with (the code for) a basic neural classifier and will have to make adaptations to that. Second, you will load existing word embeddings and perform small experiments on them.

What you have to hand in:

- A report that answers the questions. This time, you don't have to format it as a research paper, just answer the questions! Both Word and Latex are fine, as long as you submit a PDF. There's also no need to submit your code, as it will likely be very similar to the script you start out with.

**Deadline: Monday, September 23th, 09.59 AM**

## Data and Software

For this assignment, we will be using a set of Named Entity Disambiguation data, extracted from the OntoNotes corpus. It consists of single-word named entities, labelled as belonging to one of five different categories: geo-political entity (GPE), person (PERSON), organisation (ORG), date (DATE) and cardinal number (CARDINAL). The data is in the files `NE_train.txt` and `NE_dev.txt` on Brightspace.

**Note:** In addition to scikit-learn, we are using the Python deep learning library Keras in this week's assignment. If you work on your own machine, make sure to install Keras (and Tensorflow) first.

For those of you not familiar with Keras, there is an overview here, and in general the excellent documentation at <https://keras.io> should help you out.

## 2.1 – Using a neural network for named entity disambiguation

You are given five files for this assignment:

- `NE_train.txt`: a corpus of single-word named entities, and their named entity labels. Each training example is on one line, which consists of a word, followed by its label, and separated by a tab.
- `NE_dev.txt`: the development set of the corpus in similar format as the training set.
- `NE_test_input.txt`: the test set of single-word named entities. The gold standard labels are only available to us.
- `glove_filtered.json`: a Python dictionary containing the 300-dimensional word vectors, taken from here: <http://nlp.stanford.edu/projects/glove/>. Provided as a json-file, so it's quicker to read in.

**Note:** for efficiency, this set of embeddings contains only lower-cased, single words that occur in the named entity data set. You cannot use these embeddings for anything else.

- `lfd_assignment2.py`: a Python script, similar to the ones from previous assignments, that does basic neural network classification with (sort of) default parameters, and reads in the data and word embeddings. This is the file you have to experiment with! Its performance is quite bad, currently.

**Note:** make your life easier by going to "Content and Materials" for Week 3 on Brightspace and just download all the files at once using the download button.

### Live Competition

Currently, the script you're given uses a neural network with 1 layer to perform 5-class classification. For each input word, it uses the Glove embeddings as features. It reaches a performance of around 50% accuracy, which is quite bad and can surely be improved.

We will have a nice **live competition** during the lab! You have to experiment with the neural network classifier and see if you can push its performance. We have a held out test set on which we will measure your performance. The input for this test set can be downloaded from Brightspace (`NE_test_input.txt`). You get **at most 5** submissions to get a score on this test set. Your highest score will count towards the competition. We will show the current leaderboard live in class. The best performing students get a bonus on the assignment.

To submit a system, send your output file with predictions (1 prediction per line) to this Dropbox folder:

<https://www.dropbox.com/request/a5oMH7HIHp8Alhw1Y2d1>

There is no need to make an account, but you have to add your name (we will show this on the leaderboard) and your email. Please name your submission `model_{run_number}` (no extension). Dropbox automatically uses your name to rename the file (yes really), so no need to add that in the filename already.

To remind you, here are some things you can experiment with:

- Number of epochs
- Learning rate
- Batch size
- Number of layers
- Number of nodes of the hidden layers
- Activation function
- Optimizer function
- Loss function
- Addition of dropout (and the percentage)

Obviously, do not use any algorithms from previous week, any functionality like RNNs, LSTMs or Transformers. You might want to check out the Keras documentation to get up to speed.

## Questions

Of course, you can work on the assignment longer than just during the lab. You can even try to optimize your model a bit more, potentially trying different settings or architectures. For the assignment, you have to answer these questions related to neural networks:

**NN-1 (1p):** First things first: what are the majority class baseline scores? Given these scores, do you think your model reached a good performance?

**NN-2 (2p):** Please explain in detail what dropout is and how it works. Did you experiment with using it? Where did you add it in your model? Did it improve performance?

**NN-3 (2p):** How did you decide when to stop training (a single model)? Do you think this worked well? Do you think there could be better methods? Why is it important to not train for too long?

**NN-4 (2p):** With what hyperparameters did you experiment? What did you find to be the most important ones? How did you try to find the most optimal values? What other methods would there be?

**NN-5 (3p):** Please describe your final model and how exactly you arrived there. Be specific! Describe both the architecture (number of layers, dropout, activation) as well as the hyperparameter settings (learning rate, number of layers/nodes, dropout percentage). If some settings lead to improvement, please also mention specific accuracies.

**NN-6 (2p):** Look at a number of mistakes your model made on the test set (your own one). What were the most surprising errors the model made? Is there anything in particular that stands out to you? Do you think there's room for improvement?

Again, no need to structure it like a research paper, just answer the questions!